




Mapeo Objeto Relacional

May 11, 2012

Base de Datos


1



Merging Relational and Object Models

- Object-oriented models support interesting data types --- not just flat files.
 - Maps, multimedia, etc.
- The relational model supports very-high-level queries.
- Object-relational databases are an attempt to get the best of both.

2



Complex Data Types

- Motivation:
 - Permit non-atomic domains (atomic \equiv indivisible)
 - Example of non-atomic domain: set of integers, or set of tuples
 - Allows more intuitive modeling for applications with complex data
- Intuitive definition:
 - allow relations whenever we allow atomic (scalar) values
 - relations within relations
 - Retains mathematical foundation of relational model
 - Violates first normal form.



Example of a Nested Relation

- Example: library information system
- Each book has
 - title,
 - a list (array) of authors,
 - Publisher, with subfields *name* and *branch*, and
 - a set of keywords
- Non-1NF relation *books*

<i>title</i>	<i>author_array</i>	<i>publisher</i> (<i>name</i> , <i>branch</i>)	<i>keyword_set</i>
Compilers	[Smith, Jones]	(McGraw-Hill, New York)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}



Complex Types and SQL

- Extensions introduced in SQL:1999 to support complex types:
 - Collection and large object types
 - Nested relations are an example of collection types
 - Structured types
 - Nested record structures like composite attributes
 - Inheritance
 - Object orientation
 - Including object identifiers and references
- Not fully implemented in any database system currently
 - But some features are present in each of the major commercial database systems
 - Read the manual of your database system to see what it supports



User Defined Types

- A *user-defined type*, or UDT, is essentially a class definition, with a structure and methods.
- Two uses:
 1. As a *rowtype*, that is, the type of a relation.
 2. As the type of an attribute of a relation.



Structured Types and Inheritance in SQL

- Structured types (a.k.a. **user-defined types**) can be declared and used in SQL

```
create type Name as
  (firstname varchar(20),
   lastname varchar(20))
final
```

```
create type Address as
  (street varchar(20),
   city varchar(20),
   zipcode varchar(20))
not final
```

- Note: **final** and **not final** indicate whether subtypes can be created
- Structured types can be used to create tables with composite attributes

```
create table person (
  name Name,
  address Address,
  dateOfBirth date)
```
- Dot notation used to reference components: *name.firstname*



Structured Types (cont.)

- User-defined row types**

```
create type PersonType as (
  name Name,
  address Address,
  dateOfBirth date)
not final
```

- Can then create a table whose rows are a user-defined type

```
create table customer of CustomerType
```
- Alternative using **unnamed row types**.

```
create table person_t (
  name row(firstname varchar(20),
            lastname varchar(20)),
  address row(street varchar(20),
              city varchar(20),
              zipcode varchar(20)),
  dateOfBirth date)
```



Constructor Functions

- Constructor functions** are used to create values of structured types
- E.g.

```
create function Name(firstname varchar(20), lastname varchar(20))
returns Name
begin
  set self.firstname = firstname;
  set self.lastname = lastname;
end
```
- To create a value of type *Name*, we use

```
new Name('John', 'Smith')
```
- Normally used in insert statements

```
insert into Person values
  (new Name('John', 'Smith'),
   new Address('20 Main St', 'New York', '11001'),
   date '1960-8-22');
```



Type Inheritance

- Suppose that we have the following type definition for people:
create type *Person*
 (*name* **varchar**(20),
 address **varchar**(20))
- Using inheritance to define the student and teacher types
create type *Student*
 under *Person*
 (*degree* **varchar**(20),
 department **varchar**(20))
create type *Teacher*
 under *Person*
 (*salary* **integer**,
 department **varchar**(20))
- Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration



Multiple Type Inheritance

- SQL:1999 and SQL:2003 do not support multiple inheritance
- If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:
create type *Teaching Assistant*
 under *Student*, *Teacher*
- To avoid a conflict between the two occurrences of *department* we can rename them
create type *Teaching Assistant*
 under
 Student **with** (*department* **as** *student_dept*),
 Teacher **with** (*department* **as** *teacher_dept*)
- Each value must have a **most-specific type**



Array and Multiset Types in SQL

- Example of array and multiset declaration:
create type *Publisher* **as**
 (*name* **varchar**(20),
 branch **varchar**(20));
create type *Book* **as**
 (*title* **varchar**(20),
 author_array **varchar**(20) **array** [10],
 pub_date **date**,
 publisher *Publisher*,
 keyword-set **varchar**(20) **multiset**);
create table *books* **of** *Book*;



Creation of Collection Values

- Array construction
array ['Silberschatz', 'Korth', 'Sudarshan']
- Multisets
multiset ['computer', 'database', 'SQL']
- To create a tuple of the type defined by the books relation:
(**Compilers**, **array**['Smith', 'Jones'],
new Publisher (' McGraw-Hill', 'New York'),
multiset [' parsing', 'analysis'])
- To insert the preceding tuple into the relation books
insert into books
values
(**Compilers**, **array**['Smith', 'Jones'],
new Publisher (' McGraw-Hill', 'New York'),
multiset [' parsing', 'analysis']);



Unnesting

- The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.
- E.g.
select title, A as author, publisher.name as pub_name,
publisher.branch as pub_branch, K.keyword
from books as B, unnest(B.author_array) as A(author),
unnest(B.keyword_set) as K(keyword)
- Result relation *flat_books*

title	author	pub_name	pub_branch	keyword
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web



Querying Collection-Valued Attributes

- To find all books that have the word "database" as a keyword,
select title
from books
where 'database' in (unnest(keyword-set))
- We can access individual elements of an array by using indices
 - E.g.: If we know that a particular book has three authors, we could write:
select author_array[1], author_array[2], author_array[3]
from books
where title = 'Database System Concepts'
- To get a relation containing pairs of the form "title, author_name" for each book and each author of the book
select B.title, A.author
from books as B, unnest(B.author_array) as A(author)
- To retain ordering information we add a **with ordinality** clause
select B.title, A.author, A.position
from books as B, unnest(B.author_array) with ordinality as
A(author, position)

Nesting

- **Nesting** is the opposite of unnesting, creating a collection-valued attribute

- Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset

- To nest the *flat_books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch ) as  
publisher,
```

```
       collect (keyword) as keyword_set
```

```
from flat_books
```

```
groupby title, author, publisher
```

- To nest on both authors and keywords:

```
select title, collect (author) as author_set,  
       Publisher (pub_name, pub_branch) as publisher,
```

```
       collect (keyword) as keyword_set
```

```
from flat_books
```

```
group by title, publisher
```

Nesting (Cont.)

- Another approach to creating nested relations is to use subqueries in the **select** clause, starting from the 4NF relation *books4*

```
select title,
```

```
       array (select author
```

```
               from authors as A
```

```
               where A.title = B.title
```

```
               order by A.position) as author_array,
```

```
       Publisher (pub_name, pub_branch) as publisher,
```

```
       multiset (select keyword
```

```
                  from keywords as K
```

```
                  where K.title = B.title) as keyword_set
```

```
from books4 as B
```

Storing Nested Relations

- Oracle doesn't really store each nested table as a separate relation --- it just makes it look that way.
- Rather, there is one relation *R* in which all the tuples of all the nested tables for one attribute *A* are stored.
- Declare in CREATE TABLE by:
 NESTED TABLE *A* STORE AS *R*



Example: Storing Nested Tables

```
CREATE TABLE Manfs (  
  name CHAR(30),  
  addr CHAR(50),  
  beers beerTableType  
)  
NESTED TABLE beers STORE AS BeerTable;
```

19



References

- If T is a type, then **REF** T is the type of a reference to T , that is, a pointer to an object of type T .
- Often called an "object ID" in OO systems.
- Unlike object ID's, a REF is visible, although it is gibberish.

20



Object-Identity and Reference Types

- Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:
create type *Department* (
 name **varchar** (20),
 head **ref** (*Person*) **scope** *people*)
 - We can then create a table *departments* as follows
create table *departments* **of** *Department*
 - We can omit the declaration **scope** *people* from the type declaration and instead make an addition to the **create table** statement:
create table *departments* **of** *Department*
 (*head* **with options** *scope* *people*)
 - Referenced table must have an attribute that stores the identifier, called the **self-referential attribute**
create table *people* **of** *Person*
 ref is *person_id* **system generated**;



Initializing Reference-Typed Values

- To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments
values ( 'CS', null)
update departments
set head = (select p.person_id
            from people as p
            where name = 'John')
where name = 'CS'
```



Object Identifiers Using Reference Types

- **Reference type**
 - Create unique system-generated object identifiers
 - Examples:
 - REF IS SYSTEM GENERATED
 - REF IS <OID_ATTRIBUTE>
<VALUE_GENERATION_METHOD> ;



User Generated Identifiers

- The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- The table definition must specify that the reference is user generated

```
create type Person
(name varchar(20)
 address varchar(20))
ref using varchar(20)
create table people of Person
ref is person_id user generated
```
- When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address) values
('01284567', 'John', '23 Coyote Run')
```
- We can then use the identifier value when inserting a tuple into *departments*
 - Avoids need for a separate query to retrieve the identifier:

```
insert into departments
values('CS', '02184567')
```



User Generated Identifiers (Cont.)

- Can use an existing primary key value as the identifier:

```
create type Person
  (name varchar (20) primary key,
   address varchar(20))
ref from (name)
create table people of Person
ref is person_id derived
```

- When inserting a tuple for *departments*, we can then use

```
insert into departments
values(`CS`, `John`)
```



Path Expressions

- Find the names and addresses of the heads of all departments:

```
select head -> name, head -> address
from departments
```

- An expression such as "head->name" is called a **path expression**
- Path expressions help avoid explicit joins
 - If department head were not a reference, a join of *departments* with *people* would be required to get at the address
 - Makes expressing the query much easier for the user



Implementing O-R Features

- Similar to how E-R features are mapped onto relation schemas
- Subtable implementation
 - Each table stores primary key and those attributes defined in that table
- or,
- Each table stores both locally defined and inherited attributes



Presentación

- Esta presentación fue armada utilizando, además de material propio, material provisto por los siguientes autores:
- Sibberschat, Korth, Sudarshan - Database Systems Concepts, 6th Ed., Mc Graw Hill, 2010
- García Molina/Ullman/Widom - Database Systems: The Complete Book, 2nd Ed., Prentice Hall, 2009
- Elmasri/Navathe - Fundamentals of Database Systems, 6th Ed., Addison Wesley, 2011

May 11, 2012

28
